

Day 1

Virtualization

This is the process of running multiple OS's parallelly on a single piece of h/w. Here we have h/w(bare metal) on top of which we have host os and on the host os we install an application called as hypervisor. On the hypervisor we can run any no of OS's as guest OS.

The disadvantage of this approach is these applications running on the guest OS have to pass through a number of layers to access the H/W resources.

Containerization

Here we have bare metal on top of which we install the host OS and on the host OS we install an application called as Docker Engine. On the Docker engine we can run any application in the form of containers. Docker is a technology for creating these containers.

Docker achieves what is commonly called as "process isolation" i.e. all the applications (processes) have some dependency on a specific OS. This dependency is removed by Docker and we can run them on any OS as containers if we have Docker engine installed.

These containers pass through less no of layers to access the h/w resources. Also, organizations need not spend money on purchasing licenses of different OS's to maintain various applications.

Docker can be used at the stages of S/W development life cycle: Build--->Ship--->Run.

Docker comes in 2 flavours:

Docker CE (Community Edition)

Docker EE (Enterprise Edition)

Setup of Docker on Windows

1. Download Docker Desktop from <https://www.docker.com/products/docker-desktop>

2. Install it

3. Once Docker is installed we can use Power Shell to run the Docker commands

Day 2

Create an Ubuntu Linux machine using Vagrant

- 1 Download oracle virtual box from
<https://www.virtualbox.org/wiki/Downloads>
- 2 Install it
- 3 Download and install vagrant
<https://www.vagrantup.com/downloads>
- 4 Download the vagrant file and copy it into an empty folder
- 5 Open cmd prompt
- 6 Change directory to the folder where the vagrantfile is copied
cd path_of_folder
- 7 vagrant up
- 8 Username and password is:vagrant

=====

Using AWS

=====

- 1 Login in AWS account
- 2 Create a new Ubuntu 20 instance
- 3 To connect to this ubuntu instance use gitbash
<https://git-scm.com/downloads>

=====\\

Installing docker on Linux

=====

- 1 Open get.docker.com
- 2 Copy and paste the below 2 commands

```
curl -fsSL https://get.docker.com -o get-docker.sh
```

```
sh get-docker.sh
```

=====

Images and Containers

=====

A Docker image is a combination of bin/libs that are necessary for a s/w application to work. Initially all the s/w's of docker are available in the form of docker images

A running instance of an image is called as a container

=====

Docker Host: The server where docker is installed is called docker host

Docker client: This is CLI of docker which accepts the docker commands from the users and passes to a background process called docker daemon

Docker daemon: This accepts the commands coming from docker client and routes them to work on docker images or containers or the registry

Docker registry: This is the location where docker images are stored

This is of 2 type

1 Public (hub.docker.com)

2 Private: This is set up on one of our internal servers

=====
Day 3
=====

Important Docker commands

=====
Working on docker images

1 To download a docker image
docker pull image_name

2 To upload an image into docker registry
docker push image_name/image_id

3 To search for an image
docker search image_name

4 To get detailed info about an image
docker image inspect image_name/image_id

5 To delete a docker image
docker rmi image_name/image_id

6 To create a docker image from a container
docker commit container_name/container_id image_name

7 To create a docker image from a docker file
docker build -t image_name .

8 To see the list of images that are present on our docker host
docker images
or
docker image ls

9 To delete all the images
docker system prune -af

10 To see the complete history for a docker image
docker image history image_name/image_id

=====
Working on docker containers
=====

11 To see the list of all the running containers
docker container ls

12 To see the list of all the containers (running and stopped)
docker ps -a

13 To start a stopped container
docker start container_name/container_id

14 To stop a running container

```

docker stop container_name/container_id

15 To restart a container
docker restart container_name/container_id

16 To delete a stopped container
docker rm container_name/container_id

17 To delete a running container
docker rm -f container_name/container_id

18 To stop all running containers
docker stop $(docker ps -aq)

19 To delete all the stopped containers
docker rm $(docker ps -aq)

20 To delete all containers (running and stopped)
docker rm -f $(docker ps -aq)

21 To see the logs generated by a container
docker logs container_name/container_id

22 To see the ports used by the container
docker port container_name/container_id

23 To run an application in a container from outside
docker exec -it container_name/container_id command
Eg: To start a bash shell in the container
docker exec -it container_name/container_id bash

24 To go into a container
docker attach container_name/container_id

25 To create a docker container
docker run image_name/image_id
Run command options
-----
--name: Used to give a customised name to the container
-it: Used to open interactive terminal in the container
-d : Used to run a container in the background in detached mode
-e: Used to pass environment variables to a container
-v: Used to mount an external device or folder as a volume
--volumes-from: Used to share volumes between multiple containers
-p: Used for port mapping, it will link the container port with
    a port of the host machine
    eg: -p 8080:80 Here 80 is the container port(internal port) and
        8080 is the host port(external port).
-P : This will automatically perform port mapping it will link
    the container port with a host port that is greater than 30000
--link : Used to create a microservices architecture where we can link
    multiple containers
-rm: Used to remove a container on exit
--network: Used to run a container on a specific network
-c: Used to specify an upper limit on the amount of cpu that a
container
    can use
-m: Used to specify an upper limit on the memory that a container can
use

```

```

=====
==
Day 4
=====
===
Working on docker networks
=====
26 To see the list of all docker networks
    docker network ls

27 To create a network
    docker network create --driver network_type network_name

28 To get detailed info about a network
    docker network inspect network_name/network_id

29 To delete a docker network
    docker network rm network_name/network_id

30 To Connect a running container to a network
    docker network connect network_name/network_id
    container_name/container_id

31 To disconnect a running container to a network
    docker network disconnect network_name/network_id
    container_name/container_id

=====
Working on docker volumes
=====
32 To see the list of volumes
    docker volume ls

33 To create a new docker volume
    docker volume create volume_name

34 To delete a docker volume
    docker volume rm volume_name/volume_id

35 To get detailed info about a volume
    docker volume inspect volume_id/volume_name

=====
=====
=====
UseCase 1
=====
Create an nginx container in detached mode
docker run --name webserver -p 8888:80 -d nginx

To check if the nginx container is running
docker container ls

To access the nginx from browser
public_ip_dockerhost:8888
=====
===
UseCase 2
=====
Create a tomcat container and name it appserver

```

```
docker run --name appserver -p 9090:8080 -d tomee
```

To access tomcat from browser
Public_ip_of_dockerhost:9090

```
=====
=====
```

UseCase 3

```
=====
```

Create jenkins container and do automatic port mapping
docker run --name jenkins -d -P jenkins

To see the ports used by jenkins
docker port jenkins

To access the jenkins from browser
public_ip_of_dockerhost:port_no_from_previous_step

```
=====
```

UseCase 4

Start centos as a container and launch interactive terminal on it
docker run --name c1 -it centos
exit

```
=====
```

Usecase 5

Create an ubuntu container and launch interactive terminal in it
docker run --name u1 -it ubuntu
exit

```
=====
=====
```

UseCase 6

Create a mysql container and go into its bash shell
Login as mysql root user and create few tables.

1 Create a mysql container

```
docker run --name db -d -e MYSQL_ROOT_PASSWORD=intelliqit mysql:5
```

2 To open interactive bash shell in the container

```
docker exec -it db bash
```

3 To login into the db as root user

```
mysql -u root -p
```

Enter password "intelliqit"

4 To see the list of available databases

```
show databases;
```

5 To move into any of the above database

```
use db_name;
```

Eg: use sys;

6 To create emp and dept tables here

Open <https://justinsomnia.org/2009/04/the-emp-and-dept-tables-for-mysql/>

Copy the code from emp and dept tables and paste in the mysql container

7 To check if the emp and dept tables are created

```
select * from emp;
select * from dept;
```

=====

Day 5

=====

=====

Creating a multi container architecture(microservices architecture)

=====

Various types of development and testing environments can be created using docker in the following ways

- 1 --link option (depricated)
- 2 docker-compose
- 3 Docker networking

4 Python Scripts

5 Ansible

--link Option: This is a docker run command option and it is depricated.

UseCase 1

=====

Create 2 busybox containers and link them

- 1 Create a busybox container
docker run --name c1 -it busybox
- 2 To come out of the container without exit
ctrl+p,ctrl+q
- 3 Create another busybox container and link it with the c1 container
docker run --name c2 --link c1:mybusybox -it busybox
- 4 Check if c2 can ping to c1
ping c1

UseCase 2

=====

Create a mysql container and a wordpress container and link them

- 1 Create mysql:5 as a container
docker run --name db -d -e MYSQL_ROOT_PASSWORD=intelliqit mysql:5
- 2 Create a wordpress container and link it with a mysql container
docker run --name mywordpress -d -p 8888:80 --link db:mysql wordpress
- 3 To check if wordpress is linked with mysql container
docker inspect mywordpress
Search for "Links" section
- 4 To acces the wordpress from browser
public_ip_of_dockerhost:8888

UseCase 3

=====

Create a jenkins container and link with 2 tomcat containers
one for QAserver and another for prodserver

1 Create a jenkins container

```
docker run --name jenkins -d -p 5050:8080 jenkins
```

2 To access jenkins from browser

```
public_ip_of_dockerhost:5050
```

3 Create a tomcat container as qaserver and link with jenkins

```
docker run --name qaserver -d -p 6060:8080 --link jenkins:myjenkins  
tomcat
```

4 Create another tomcat container as prodserver and link with jenkins

```
docker run --name prodserver -d -p 7070:8080 --link jenkins:myjenkins  
tomcat
```

=====

Day 6

=====

Use Case

=====

Setup a postgres db and link with adminer container to access db from
browser

1 Create a postgres db container

```
docker run --name db -d -e POSTGRES_PASSWORD=intelliqit -e  
POSTGRES_USER=myuser  
postgres -e POSTGRES_DB=mydb
```

2 Create an adminer container and link with postgres

```
docker run --name myadminer -d -p 9999:8080 --link db:postgres adminer
```

3 To access from level of browser

```
public_ip_of_dockerhost:9999
```

=====

UseCase

=====

Setup LAMP architecture where a mysql container can be linked
with an apache and php container

1 Create a mysql container

```
docker run --name db -d -e MYSQL_ROOT_PASSWORD=intelliqit mysql
```

2 Create an apache and link with mysql container

```
docker run --name apache -d -p 9090:80 --link db:mysql httpd
```

3 Create a php container and link with apache and mysql containers

```
docker run --name php -d --link db:mysql --link apache:httpd php:7.2-  
apache
```

4 To check if php container is linked with mysql and apache container

```
docker inspect php
```

Search for "Links" section


```
=====
Day 7
=====
=====
=====
```

UseCase

```
=====
```

Create a testing environment where a selenium hub container should be linked with 2 node containers one with chrome installed and other with firefox installed. The testers should be able to run the cross browser, cross platform automation test scripts

- 1 Create a selenium hub image
docker run --name hub -d -p 4444:4444 selenium/hub
- 2 Create a chrome node and link with the hub container
docker run --name chrome -d -p 5901:5900 --link hub:selenium selenium/node-chrome-debug
- 3 Create a firefox node and link with the hub container
docker run --name firefox -d -p 5902:5900 --link hub:selenium selenium/node-firefox-debug
- 4 Check if all 3 containers are running
docker container ls
- 5 The above 2 containers are GUI containers and to access the GUI of these containers
 - a) Install VNC viewer from <https://www.realvnc.com/en/connect/download/viewer/>
 - b) Open vnc viewer
 - c) public_ip_of_dockerhost:5901 and 5902
 - d) Click on continue--->Enter password:secret

```
=====
=====
```

Docker Compose

```
=====
```

This is used for creating a multi container architecture which is reusable
Docker compose uses yaml files

Install docker-compose

```
=====
```

- 1 Open <https://docs.docker.com/compose/install/>
- 2 Click on Linux tab
- 3 Copy and paste the commands
- 4 To check the version of docker compose
docker-compose --version

```
=====
=====
```

Day 8

```

=====
=====
Docker Compose
=====
This is used for creating a multi container architecture which is
reusable
Docker compose uses yml files

Install docker-compose
=====
1 Open https://docs.docker.com/compose/install/
2 Click on Linux tab
3 Copy and paste the commands

4 To check the version of docker compose
  docker-compose --version

=====
=
UseCase
Create a docker compose file to setup a mysql container linked with
a database container

vim docker-compose.yml

---
version: '3.8'

services:
  mydb:
    image: mysql:5
    environment:
      MYSQL_ROOT_PASSWORD: intelligit

  mywordpress:
    image: wordpress
    image:
    ports:
      - "8888:80"
    links:
      - mydb:mysql
  ...

To setup the containers from the above file
docker-compose up -d

To stop the containers
docker-compose stop

To stop and delete the containers
docker-compose down

=====
UseCase
Create a docker compose file to setup ci-cd environment
where a jenkins is linked with 2 tomcat containers

vim docker-compose.yml
version: '3.8'

```

```
services:
  myjenkins:
    image: jenkins/jenkins
    ports:
      - 5050:8080
    container_name: myjenkins
```

```
qaserver:
  image: tomcat
  ports:
    - 6060:8080
  links:
    - myjenkins:jenkins
  container_name: qaserver
```

```
prodserver:
  image: tomcat
  ports:
    - 7070:8080
  links:
    - myjenkins:jenkins
  container_name: prodserver
```

...

=====

==

UseCase

=====

Create a docker compsoe file to setup the LAMP architecture

version: '3.8'

```
services:
  mydb:
    image: mysql
    environment:
      MYSQL_ROOT_PASSWORD: intelliqit
    container_name: mydb
```

```
apache:
  image: httpd
  ports:
    - 9090:80
  container_name: apache
  links:
    - mydb:mysql
```

```
php:
  image: php:7.2-apache
  links:
    - mydb:mysql
    - apache:httpd
  container_name: php
```

...

=====

UseCase

Create a docker compose file to setup start mysql and link with adminer application

vim docker-compose.yml

```

---
version: '3.8'

services:
  mydb:
    image: mysql
    environment:
      MYSQL_ROOT_PASSWORD: intelliqit
    container_name: mydb

  myadminer:
    image: adminer
    ports:
      - 8888:8080
    container_name: myadminer
...

```

Day 8

UseCase

Create a docker compose file to setup the selenium testing environment where a selenium hub container is linked with 2 node containers one with chrome and other with firefox

vim docker-compose.yml

```

---
version: '3.8'

services:
  hub:
    image: selenium/hub
    ports:
      - 4444:4444
    container_name: hub

  chrome:
    image: selenium/node-chrome-debug
    ports:
      - 5901:5900
    links:
      - hub:selenium
    container_name: chrome

  firefox:
    image: selenium/node-firefox-debug
    ports:
      - 5902:5900
    links:
      - hub:selenium
    container_name: firefox
...

```

To setup the above architecture
docker-compose up -d

To check the running containers

docker container ls

To delete the containers

docker-compose down

=====

Docker Volumes

=====

Containers are ephemeral(temporary) but the data processed by the containers should be persistent. Once a container is deleted all the data of the container is lost

To preserve the data even if the container is deleted we can use volumes

Volumes are classified into 3 types

=====

- 1 Simple docker volume
- 2 Sharable docker volumes
- 3 Docker volume containers

Simple Docker volumes

=====

These volumes are used only for preserving the data on the host machine even if the containers are deleted

Use Case

=====

Create a directory /data and mount it as a volume on an ubuntu container
Create some files in the mounted volumes and check if the files are preserved on the host machine even after the container is deleted

- 1 Create /data directory
mkdir /data
- 2 Create an ubuntu container and mount the above directory as volume
docker run --name u1 -it -v /data ubuntu
In the container u1 go into /data directory and create some files
cd /data
touch file1 file2 file3
exit
- 3 Identify the location where the mounted data is preserved
docker inspect u1
Search for "Mounts" section and copy the "Source" path
- 4 Delete the container
docker rm -f u1
- 5 Check if the data is still present
cd "source_path_from_step3"
ls

=====

Sharable Docker volumes

=====

These volumes are sharable between multiple containers

Create 3 centos containers c1, c2, c3.

Mount /data as a volume on c1 container, c2 should use the volume used by c1 and c3 should use the volume used by c2

- 1 Create a centos container c1 and mount /data
docker run --name c1 -it -v /data centos
- 2 Go into the data folder create files in data folder
cd data
touch f1 f2
- 3 Come out of the container without exit
ctrl+p,ctrl+q
- 4 Create another centos container c2 and it should use the volumes used by c1
docker run --name c2 -it --volumes-from c1 centos
- 5 In the c2 container go into data folder and create some file
cd data
touch f3 f4
- 6 Come out of the container without exit
ctrl+p,ctrl+q
- 7 Create another centos container c3 and it should use the volume used by c2
docker run --name c3 -it --volumes-from c2 centos
- 8 In the c3 container go into data folder and create some file
cd data
touch f5 f6
- 9 Come out of the container without exit
ctrl+p,ctrl+q
- 10 Go into any of the 3 containers and we will see all the files
docker attach c1
cd /data
ls
exit
- 12 Identify the location where the mounted data is stored
docker inspect c1
Search for "Mounts" section and copy the "Source" path
- 13 Delete all containers
docker rm -f c1 c2 c3
- 14 Check if the files are still present
cd "source_path_from"step12"

=====
Day 9
=====

Docker volume containers

These volumes are bidirectional i.e. the changes done on host will be reflected into container and changes done by container will be reflected to host machine

- 1 Create a volume

```
docker volume create myvolume
```

2 To check the location where the mounted the volume works
`docker volume inspect myvolume`

3 Copy the path shown in "MountPoint" and cd to that Path
`cd "MountPoint"`

4 Create few files here
`touch file1 file2`

5 Create a centos container and mount the above volume into the tmp folder
`docker run --name c1 -it -v myvolume:/tmp centos`

6 Change to tmp folder and check for the files
`cd /tmp`
`ls`
If we create any files here they will be reflected to host machine
And these files will be present on the host even after deleting the container.

```
=====
```

```
==
```

UseCase

```
=====
```

Create a volume "newvolume" and create tomcat-users.xml file in it
Create a tomcat container and mount the above volume into it
Copy the tomcat-users.xml files to the required location

1 Create a volume
`docker volume create newvolume`

2 Identify the mount location
`docker volume inspect newvolume`
Copy the "MountPoint" path

3 Move to this path
`cd "MountPoint path"`

4 Create a file called tomcat-users.xml
`cat > tomcat-users.xml`
`<tomcat-users>`
 `<user username="intelliqit" password="intelliqit" roles="manager-script"/>`
`</tomcat-users>`

5 Create a tomcat container and mount the above volume
`docker run --name webserver -d -P -v newvolume:/tmp tomcat`

6 Go into bash shell of the tomcat container
`docker exec -it webserver bash`

7 Move the tomcat-users.xml file into conf folder
`mv /tmp/tomcat-users.xml conf/`

```
=====
```

Creating customsied docker images

```
=====
```

This can be done in 2 ways

- 1 Using docker commit command
- 2 Using dockerfile

Using the docker commit command

=====

UseCase

=====

Create an ubuntu container and install some s/w's in it
Save this container as an image and later create a new container
from the newly created image. We will find all the s/w's that we
installed.

- 1 Create an ubuntu container
docker run --name u1 -it ubuntu
- 2 In the container update the apt repo and install s/w's
apt-get update
apt-get install -y git
- 3 Check if git is installed or not
git --version
exit
- 4 Save the customised container as an image
docker commit u1 myubuntu
- 5 Check if the new image is created or not
docker images
- 6 Delete the previously create ubuntu container
docker rm -f u1
- 7 Create an new container from the above created image
docker run --name u1 -it myubuntu
- 8 Check for git
git --version

=====

Dockerfile

=====

Dockerfile uses predefined keyword to create customised
docker images.

Important keyword in dockerfile

=====

FROM : This is used to specify the base image from where a
customised docker image has to be created

MAINTAINER : This represents the name of the organization or the
author that has created this dockerfile

RUN :Used to run linux commands in the container
Generally it used to do s/w installation or
running scripts

USER : This is used to specify who should be the default user

to login into the container

COPY : Used to copy files from host to the customised image that we are creating

ADD : This is similar to copy where it can copy files from host to image but ADD can also download files from some remote server

EXPOSE : Used to specify what port should be used by the container

VOLUME : Used for automatic volume mounting ie we will have a volume mounted automatically when the container start

WORKDIR : Used to specify the default working directory of the container

ENV : This is used to specify what environment variables should be used

CMD : Used to run the default process of the container from outside

ENTRYPOINT : This is also used to run the default process of the container

LABEL: Used to store data about the docker image in key value pairs

SHELL : Used to specify what shell should be by default used by the image

=====

UseCase

=====

Create a dockerfile to use nginx as base image and specify the maintainer as intelliqt

1 Create docker file
vim dockerfile

FROM nginx
MAINTAINER intelliqt

2 To create an image from this file
docker build -t mynginx .

3 Check if the image is created or not
docker images

Day 10

=====

UseCase

=====

Create a dockerfile from ubuntu base image and install git in it

1 Create dockerfile
vim dockerfile

```
FROM ubuntu
MAINTAINER intelliqit
RUN apt-get update
RUN apt-get install -y git
```

2 Create an image from the above file
docker build -t myubuntu .

3 Check if the new image is created
docker images

4 Create a container from the new image and it should have git installed
docker run --name ul -it myubuntu
git --version

=====

Cache Busting

=====

When we create an image from a dockerfile docker stores all the executed instructions in its cache. Next time if we edit the same docker file and add few new instructions and build an image out of it docker will not execute the previously executed statements. Instead it will read them from the cache. This is a time saving mechanism. The disadvantage is if the docker file is edited with a huge time gap then we might end up installing s/w's that are outdated.

Eg:

```
FROM ubuntu
RUN apt-get update
RUN apt-get install -y git
RUN apt-get install -y tree
```

If we build an image from the above dockerfile docker saves all these instructions in the dockercache and if we add the below statement

```
RUN DEBIAN_FRONTEND=noninteractive apt-get -yq install maven
```

only this latest statement will be executed

To avoid this problem and make docker execute all the instructions once more time without reading from cache we use "cache busting"
docker build --no-cache -t myubuntu .

=====

Day 11

=====

Create a shell script to install multiple s/w's and copy this into the docker image and execute it at the time of creating the image

1 Create the shell script

```
vim script.sh
apt-get update
for x in tree git
do
    apt-get install -y $x
done
```

2 Give execute permissions on that file
chmod u+x script.sh

```

3 Create the dockerfile
  vim dockerfile
  FROM ubuntu
  MAINTIANER intelliqit
  COPY ./script.sh /
  RUN ./script.sh

4 Create an image from the dockerfile
  docker build -t myubuntu .

5 Create a container from the above image
  docker run --name ul -it myubuntu

6 Check if the script.sh is present in / and also see if tree and git are
  installed
  ls /
  git --version
  tree

```

```

=====
===

```

UseCase

```

=====

```

Create a dockerfile from ubuntu base image and install ansible in it

```

1 Create a dockerfile
  FROM ubuntu
  MAINTAINER intelliqit
  RUN apt-get update
  RUN apt-get install -y software-properties-common
  RUN apt-get install -y ansible

2 Create an image from the above file
  docker build -t ansible .

3 Create a container from the above image and check if ansible is present
  docker run --name a1 -it ansible
  ansible --version

```

```

=====
=====

```

Create a dockerfile from ubuntu base image and make /data as the default volume

```

1 Create a dockerfile
  vim dockerfile
  FROM ubuntu
  MAINTAINER intelliqit
  VOLUME /data

2 Create an image from the above dockerfile
  docker build -t myubuntu .

3 Create a container from the above image and check for the volume
  docker run --name ul -it myubuntu
  ls (we should see the data folder)

4 Go into the data folder and create some files

```

```
cd data
touch file1 file2
exit
```

5 Check for the mount section and copy the source path

6 Delete the container
docker rm -f u1

7 Check if the files are still present
cd "source path"

=====

Create a dcoker file from nginx base image and expose 80 port

1 vim dockerfile
FROM nginx
MAINTAINER intelliqit
EXPOSE 90

2 Create an image
docker build -t mynginx .

3 Create a container from above image
docker run --name n1 -d -P mynginx

4 Check for the ports exposed
docker port n1

=====

Create a dockerfile from ubuntu base image and downlaod jenkins.war into it

1 Create a dockerfile
vim dockerfile
FROM ubuntu
MAINTIANER intelliqit
ADD https://get.jenkins.io/war-stable/2.263.4/jenkins.war /

2 Create an image from the above dockerfile
docker build -t myubuntu .

4 Create a container from this image
docker run --name u1 -it myubuntu

5 Check if jenkins.war is present
ls

=====

=====

Day 12

=====

===

Create a dockerfile from jenkins base image and make root as the deafulst user

1 vim dockerfile
FROM jenkins/jenkins
MAINTAINER intelliqit
USER root

- 2 Create an image from the above dockerfile
docker build -t myjenkins .
- 3 Create a container
docker run --name j1 -d -P myjenkins
- 4 Go into the bash shell and check if the user is root
docker exec -it j1 bash
whoami

=====

Create a dcokerfile from ubuntu base image and install java in it,docwnload jenkins.war and make "java -jar jenkins.war" as the default process of the container

- 1 vim dockerfile
FROM ubuntu
MAINTAINER intelliqit
RUN apt-get update
RUN apt-get install -y openjdk-8-jdk
ADD https://get.jenkins.io/war-stable/2.263.4/jenkins.war /
ENTRYPOINT ["java","-jar","jenkins.war"]
- 2 Create an image from the above file
docker build -t myubuntu .
- 3 Create a container from the above image and we will see that it behaves like a jenkins container
docker run --name u1 -it myubuntu
- 4 Check the default process that is running
docker container ls

=====

=====

Day 13

=====

=====

UseCase

=====

Create a dockerfile from ubuntu base image and make it behave like nginx

- 1 Create a dockerfile
vim dockerfile
FROM ubuntu
MAINTAINER intelliqit
RUN apt-get update
RUN apt-get install -y nginx
ENTRYPOINT ["/usr/sbin/nginx","-g","daemon off;"]
EXPOSE 80
- 2 Create an image from the above dockerfile
docker build -t myubuntu .
- 3 Create a container from the above image and it will work like nginx
docker run --name n1 -d -P myubuntu

4 Check the ports used by nginx
docker container ls

5 To access nginx from browser
public_ip_of_dockerhost:port_no_captured_from_step4

=====

====
UseCase

Create a dockerfile fromj centos base image and install
httpd in it make httpd as the default process

1 Create the index.html

```
<html>
    <body>
        <h1>Welcome to IntelliQIT</h1>
    </body>
</html>
```

2 Create the dockerfile

```
vim dockerfile
FROM centos
MAINTAINER intelliqit
RUN yum -y update
RUN yum -y install httpd
COPY index.html /var/www/html
ENTRYPOINT ["/usr/sbin/httpd","-D","FOREGROUND"]
EXPOSE 80
```

3 Create an image from the above dockerfile
docker build -t mycentos .

4 Create a container from the above image
docker run --name c1 -d -P mycentos

5 Check the ports used by container
docker container ls

6 To access the from browser
public_ip_of_dockerhost:port_from_step5

=====

CMD and ENTRYPOINT

Bothe of them are used to specify the default process that should be
triggered when the container starts but the CMD instruction can be
overridden with some other process passed at the docker run command

Eg:

```
FROM ubuntu
RUN apt-get update
RUN apt-get install -y nginx
CMD ["/usr/sbin/nginx","-g","daemon off;"]
EXPOSE 80
```

Though the default process is to trigger nginx we can bypass that
and make it work on some other process

```
docker build -t myubuntu .
Create a container
docker run --name u1 -it -d myubuntu
Here if we inspect the default process we will see that
nginx as the default process
docker container ls
```

on the otherhand we can modify that default process to something else
docker run --name u1 -d -P myubuntu ls -la
Now if we do "docker container ls" we will see the default process
to be "ls -la"

```
=====
===
Docker Networking
=====
Docker uses 4 types os networks
1 Bridge: This is the deaefault network of docker when contianers are
           running on a single docker host

2 Host: This is used when we want to run a single container on a
dockerhost
           and this contianer communicates only with the host machine

3 Null: This is used for creating isolated containers ie these containers
        cannot communicate with th host machine or with other containers

4 Overlay: This is used when containers are running in a distributed
environment
           on multiple linux servers
```

```
UseCase
=====
Create 2 bridge networks intelliq1 and intelliq2
Create 2 busybox containers c1,c2 and c3
c1 and c2 should run on intelliq1 network and shoul ping each other
c3 should run on intelliq2 network and it should not be able to ping c1
or c2
Now put c2 on intelliq2 network,since c2 is on both intelliq1 and
intelliq2
networks it should be able to ping to both c1 and c3
but c1 and c3 should not ping each other directly

1 Create 2 bridge networks
  docker network create --driver bridge intelliq1
  docker network create --driver bridge intelliq2

2 Check the list of available networks
  docker network ls

3 Create a busybox container c1 on intelliqi1 network
  docker run --name c1 -it --network intelliq1 busybox
  Come out of the c1 container without exit ctrl+p,ctrl+q

4 Identify the ipaddress of c1
  docker inspect c1
```

- 5 Create another busybox container c2 on intelliq1 network
 docker run --name c2 -it --network intelliq1 busybox
 ping ipaddress_of_c1 (It will ping)
 Come out of the c2 container without exit ctrl+p,ctrl+q
- 6 Identify the ipaddress of c2
 docker inspect c2
- 7 Create another busybox container c3 on intelliq2 network
 docker run --name c3 -it --network intelliq2 busybox
 ping ipaddress_of_c1 (It should not ping)
 ping ipaddress_of_c2 (It should not ping)
 Come out of the c3 container without exit ctrl+p,ctrl+q
- 8 Identify the ipaddress of c3
 docker inspect c3
- 9 Now attach intelliq2 network to c2 container
 docker network connect intelliq2 c2
- 10 Since c2 is now on both intelliq1 and intelliq2 networks it should
 ping
 to both c1 and c3 containers
 docker attach c2
 ping ipaddress_of_c1 (It should ping)
 ping ipaddress_of_c3 (It should ping)
 Come out of the c2 container without exit ctrl+p,ctrl+q
- 11 But c1 and c3 should not ping each other
 docker attach c3
 ping ipaddress_of_c1 (It should not ping)

=====
 ==

Day 14

=====
 ==

=====
 =====

=====
 UseCase

=====

Create a custom bridge network and create a docker compose file
 to start postgres and adminer container on the above created
 network

- 1 Create a custom bridge network
 docker network create --driver bridge --subnet 10.0.0.0/24 intelliq1

- 2 Create a docker compose file
 vim docker-compose.yml

version: '3.8'

services:

db:

image: postgres

environment:

POSTGRES_PASSWORD: intelliq1

POSTGRES_USER: myuser

POSTGRES_DB: mydb


```
adminer:
  image: adminer
  ports:
    - 8888:8080
```

```
networks:
  default:
    external:
      name: intelliqit
  ...
```

3 To create the containers
docker-compose up -d

4 To see if adminer and postgres containers are created
docker container ls

5 To check if they are running on intelliqit network
docker inspect container_id_from_Step4

=====

UseCase

Create a docker compose file which starts jenkins and 2 tomee as containers also this compose file should create 2 networks abc,xyz.On abc network it should run the jenkins container and xyz network it should run the 2 tomee containers

vim docker-compose.yml

version: '3.8'

```
services:
  myjenkins:
    image: jenkins/jenkins
    ports:
      - 5050:8080
    networks:
      - abc
```

```
qaserver:
  image: tomee
  ports:
    - 6060:8080
  networks:
    - xyz
```

```
prodserver:
  image: tomee
  ports:
    - 7070:8080
  networks:
    - xyz
```

```
networks:
  abc: {}
  xyz: {}
```


UseCase

=====

Create a docker compose file to start mysql and wordpress as container
it should also create 2 volumes one for wordpress and other for mysql

vim docker-compose.yml

version: '3.8'

services:

db:

image: mysql:5

environment:

MYSQL_ROOT_PASSWORD: intelliqit

volumes:

- db:/var/lib/mysql

wordpress:

image: wordpress

ports:

- 9090:80

volumes:

- wordpress:/var/www/html

volumes:

db:

wordpress:

To create containers from the above file

docker-compose up -d

Check if 2 new volumes are created

docker volume ls

=====

UseCase

=====

Create a dockerfile to create a customized jenkins image
and this dockerfile should be built as image from the docker composefile

vim dockerfile

FROM jenkins/jenkins

MAINTAINER intelliqit

USER root

RUN apt-get update

RUN apt-get install -y git

vim docker-compose.yml

version: '3.8'

services:

jenkins:

build: .

```
mytomcat:
  image: tomee
  ports:
    - 8080:8080
...
```

To create containers from the above file
docker-compose up -d

Check if a new jenkins image has be created
docker images

```
=====
Day 15
=====
=
=====
====
Container Orchestration
=====
This is the process of handling docker containers running
on multiple linux servers in a distributed environment
```

Advantages
=====

- 1 Load Balancing
- 2 Scalling
- 3 Rolling update
- 4 High Availability and Disaster recovery(DR)

LoadBalancing
=====

Each container is capable of sustaining a specific user load
We can increase this capacity by running the same application
on multiple containers(replicas)

Scalling
=====

We should be able to increase or decrease the number of containers
on which our applications are running without the end user
experiencing any downtime.

Rolling update
=====

Application running in a live environment should be upgraded or
downgraded to a different version without the end user having any
downtime

Disaster Recovery
=====

In case of network failuers or server crashes still the container
orchestration tools maintain the desired count of containers
and thereby provide the same service to the end user

Popular container orchestration tools
=====

- 1 Docker Swarm
- 2 Kubernetes
- 3 OpenShift

4 Mesos

```
=====
==
```

Setup of Docker Swarm

```
=====
```

- 1 Create 3 AWS ubuntu instances
- 2 Name them as Manager,Worker1,Worker2
- 3 Install docker on all of them
- 4 Change the hostname
vim /etc/hostname
Delete the content and replace it with Manager or Worker1 or Worker2
- 5 Restart
init 6
- 6 To initialise the docker swarm
Connect to Manager AWS instance
docker swarm init
This command will create a docker swarm and it will also generate a tokenid
- 7 Copy and paste the token id in Worker1 and Worker2

```
=====
=====
```

Ports used by docker swarm

```
=====
```

TCP port 2376 for secure Docker client communication. This port is required for Docker Machine to work. Docker Machine is used to orchestrate Docker hosts.

TCP port 2377. This port is used for communication between the nodes of a Docker Swarm or cluster. It only needs to be opened on manager nodes.

TCP and UDP port 7946 for communication among nodes (container network discovery).

UDP port 4789 for overlay network traffic (container ingress networking).

```
=====
=====
```

Load Balancing:

Each docker containers has a capability to sustain a specific user load.To increase this capability we can increase the number of replicas(containers) on which a service can run

UseCase

```
-----
```

Create nginx with 5 replicas and check where these replicas are running

- 1 Create nginx with 5 replicas
docker service create --name webserver -p 8888:80 --replicas 5 nginx
- 2 To check the services running in swarm
docker service ls
- 3 To check where these replicas are running
docker service ps webserver
- 4 To access the ngonx from browser

public_ip_of_manager/worker1/worker2:8888

5 To delete the service with all replicas
docker service rm webserver

=====

UseCase

=====

Create mysql with 3 replicas and also pass the necessary environment variables

1 docker service create --name db --replicas 3
-e MYSQL_ROOT_PASSWORD=intelliqit mysql:5

2 To check if 3 replicas of mysql are running
docker service ps db

=====

=

Day 16

=====

=

Scalling

=====

This is the process of increasing the number of replicas or decreasing the replicas count based on requirement without the end user experiencing any down time.

UseCase

=====

Create tomcat with 4 replicas and scale it to 8 and scale it down to 2

1 Create tomcat with 4 replicas
docker service create --name appserver -p 9090:8080 --replicas 4 tomcat

2 Check if 4 replicas are running
docker service ps appserver

3 Increase the replicas count to 8
docker service scale appserver=8

4 Check if 8 replicas are running
docker service ps appserver

5 Decrease the replicas count to 2
docker service scale appserver=2

6 Check if 2 replicas are running
docker service ps appserver

=====

Rolling updates

=====

Services running in docker swarm should be updated from once version to other without the end user downtime

UseCase

=====

Create redis:3 with 5 replicas and later update it to redis:4

also rollback to redis:3

1 Create redis:3 with 5 replicas

docker service create --name myredis --replicas 5 redis:3

2 Check if all 5 replicas of redis:3 are running

docker service ps myredis

3 Perform a rolling update from redis:3 to redis:4

docker service update --image redis:4 myredis

4 Check redis:3 replicas are shut down and in its place redis:4 replicas are running

docker service ps myredis

5 Roll back from redis:4 to redis:3

docker service update --rollback myredis

6 Check if redis:4 replicas are shut down and in its place redis:3 is running

docker service ps myredis

=====

To remove a worker from swarm cluster

docker node update --availability drain Worker1

To make this worker rejoin the swarm

docker node update --availability active Worker1

To make worker2 leave the swarm

Connect to worker2 using git bash

docker swarm leave

To make manager leave the swarm

docker swarm leave --force

To generate the tokenid for a machine to join swarm as worker

docker swarm join-token worker

To generate the tokenid for a machine to join swarm as manager

docker swarm join-token manager

To promote Worker1 as a manager

docker node promote Worker1

To demote "Worker1" back to a worker status

docker node demote Worker1

=====

FailOver Scenarios of Workers

=====

Create httpd with 6 replicas and delete one replica running on the manager

Check if all 6 replicas are still running

Drain Worker1 from the docker swarm and check if all 6 replicas are running
on Manager and Worker2, make Worker1 rejoin the swarm

Make Worker2 leave the swarm and check if all the 6 replicas are running on Manager and Worker1

- 1 Create httpd with 6 replicas
docker service create --name webserver -p 9090:80 --replicas 6 httpd
- 2 Check the replicas running on Manager
docker service ps webserver | grep Manager
- 3 Check the container id
docker container ls
- 4 Delete a replica
docker rm -f container_id_from_step3
- 5 Check if all 6 replicas are running
docker service ps webserver
- 6 Drain Worker1 from the swarm
docker node update --availability drain Worker1
- 7 Check if all 6 replicas are still running on Manager and Worker2
docker service ps webserver
- 8 Make Worker1 rejoin the swarm
docker node update --availability active Worker1
- 9 Make Worker2 leave the swarm
Connect to Worker2 using git bash
docker swarm leave
Connect to Manager
- 10 Check if all 6 replicas are still running
docker service ps webserver

=====

Day 17

=====

FailOver Scenarios of Managers

=====

If a worker instance crashes all the replicas running on that worker will be moved to the Manager or the other workers.
If the Manager itself crashes the swarm becomes headless
ie we cannot perform container orchestration activities in this swarm cluster

To avoid this we should maintain multiple managers
Manager nodes have the status as Leader or Reachable

If one manager node goes down other manager becomes the Leader
Quorum is responsible for doing this activity and it uses a RAFT algorithm for handling the failovers of managers. Quorum also is responsible for maintaining the min number of manager

Min count of manager required for docker swarm should be always more than half of the total count of Managers

Total Manager Count	-	Min Manager Required
1	-	1
2	-	2
3	-	2
4	-	3
5	-	3
6	-	4
7	-	4

```
=====
==
=====
=====
```

Overlay Networking

```
=====
```

This is the default network used by swarm
and this network performs network load balancing
ie even if a service is running on a specific worker we can
access it from other slave

```
=====
====
UseCase
=====
```

Create 2 overlay networks intelliqit1 and intelliqit2
Create httpd with 5 replicas on intelliqit1 network
Create tomcat with 5 replicas on default overlay "ingres" network
and later perform rolling network update to intelliqit2 network

- 1 Create 2 overlay networks

```
docker network create --driver overlay intelliqit1
docker network create --driver overlay intelliqit2
```
- 2 Check if 2 overlay networks are created

```
docker network ls
```
- 3 Create httpd with 5 replicas on intelliqit1 network

```
docker service create --name webserver -p 8888:80 --replicas 5
                                --network intelliqit1 httpd
```
- 4 To check if httpd is running on intelliqit1 network

```
docker service inspect webserver
```

This command will generate the output in JSON format
To see the above output in normal text format

```
docker service inspect webserver --pretty
```
- 5 Create tomcat with 5 replicas on the default ingres network

```
docker service create --name appserver -p 9999:8080 --replicas 5 tomcat
```
- 6 Perform a rolling network update from ingres to intelliqit2 network

```
docker service update --network-add intelliqit2 appserver
```


7 Check if tomcat is now running on intelligit2 network
docker service inspect appserver --pretty

Note: To remove from intelligit2 network
docker service update --network-rm intelligit2 appserver

```
=====
=====
Day 18
=====
=====
=====
```

Overlay Networking

=====

This is the default network used by swarm
and this network performs network load balancing
ie even if a service is running on a specific worker we can
access it from other slave

UseCase

=====

Start nginx with 2 replicas and check if we can access it from
browser from manager and all workers

- 1 Create nginx
docker service create --name webserver -p 8888:80 --replicas 2 nginx
- 2 Check where these 2 replicas are running
docker service ps webserver
These replicas will be running on only 2 nodes and we will have a third
node where it is not running
- 3 Check if we can access nginx from the third node where it is not
present
public_ip_of_thirdnode:8888

```
=====
=====
```

Docker Stack

=====

docker compose + docker swarm = docker stack
docker compose + kubernetes = kompose

Docker compose when implemented at the level of docker swarm
it is called docker stack. Using docker stack we can create an orchestration
a micro services architecture at the level of production servers

- 1 To create a stack from a compose file
docker stack deploy -c compose_filename stack_name
- 2 To see the list of stacks created
docker stack ls
- 3 To see on which nodes the stack services are running
docker stack ps stack_name
- 4 To delete a stack
docker stack rm stack_name

=====

UseCase

=====

Create a docker stack file to start 3 replicas of wordpress
and one replica of mysql

vim stack1.yml

version: '3.8'

services:

db:

image: "mysql:5"

environment:

MYSQL_ROOT_PASSWORD: intelliqit

wordpress:

image: wordpress

ports:

- "8989:80"

deploy:

replicas: 3

To start the stack file

docker stack deploy -c stack1.yml mywordpress

To see the services running

docker service ls

To check where the services are running

docker stack ps mywordpress

To delete the stack

docker stack rm mywordpress

=====

=====

UseCase

=====

Create a stack file to setup CI-cd architecture where a jenkins
container is linked with tomcats for qa and prod environments
The jenkins containers should run only on Manager
the qaserver tomcat should run only on Worker1 and prodserver
tomcat should run only on worker2

vim stack2.yml

version: '3.8'

services:

myjenkins:

image: jenkins/jenkins

ports:

- 5050:8080

deploy:

replicas: 2

placement:

constraints:

- node.hostname == Manager

```
qaserver:
  image: tomcat
  ports:
    - 6060:8080
  deploy:
    replicas: 3
    placement:
      constraints:
        - node.hostname == Worker1
```

```
prodserver:
  image: tomcat
  ports:
    - 7070:8080
  deploy:
    replicas: 4
    placement:
      constraints:
        - node.hostname == Worker2
```

...

To start the services
docker deploy -c stack2.yml ci-cd

To check the replicas
docker stack ps ci-cd
=====

=====

Day 19

=====

UseCase
Create a stack file to setup the selenium hub and nodes architecture
but also specify a upper limit on the h/w

```
vim stack3.yml
---
version: '3.8'
```

```
services:
  hub:
    image: selenium/hub
    ports:
      - 4444:4444
    deploy:
      replicas: 2
      resources:
        limits:
          cpus: "0.1"
          memory: "300M"
```

```
chrome:
  image: selenium/node-chrome-debug
  ports:
    - 5901:5900
  deploy:
    replicas: 3
    resources:
```

```
limits:
  cpus: "0.01"
  memory: "100M"
```

```
firefox:
  image: selenium/node-firefox-debug
  ports:
    - 5902:5900
  deploy:
    replicas: 3
  resources:
    limits:
      cpus: "0.01"
      memory: "100M"
```

=====

Docker secrets

=====

This is a feature of docker swarm using which we can pass secret data to the services running in swarm cluster

These secrets are created on the host machine and they will be available from all the replicas in the swarm cluster

- 1 Create a docker secret
echo " Hello Intelliqit" | docker secret create mysecret -
- 2 Create a redis db with 5 replicas and mount the secret
docker service create --name myredis --replicas 5 --secret mysecret redis
- 3 Capture one of the replica container id
docker container ls
- 4 Check if the secret data is available
docker exec -it container_id cat /run/secrets/mysecret

=====

=====

Create 3 secrets for postgres user,password and db and pass them to the stack file

- 1 Create secrets
echo "intelliqit" | docker secret create pg_password -
echo "myuser" | docker secret create pg_user -
echo "mydb" | docker secret create pg_db -

- 2 Check if the secrets are created
docker secret ls

- 3 Create the docker stack file to work on these secrets
vim stack6.yml

version: '3.1'

services:

```
db:
  image: postgres
  environment:
    POSTGRES_PASSWORD_FILE: /run/secrets/pg_password
    POSTGRES_USER_FILE: /run/secrets/pg_user
    POSTGRES_DB_FILE: /run/secrets/pg_db
```

```

secrets:
  - pg_password
  - pg_user
  - pg_db

adminer:
  image: adminer
  restart: always
  ports:
    - 8080:8080
  deploy:
    replicas: 2

secrets:
  pg_password:
    external: true
  pg_user:
    external: true
  pg_db:
    external: true

...
=====
Working on docker registry
=====
This is the location where the docker images are saved
This is of 2 types
1 Public registry
2 Private registry

UseCase
Create a customised ubuntu image and upload into the public registry

1 Signup into hub.docker.com

2 Create a customised ubuntu image
  a) Create a centos container and install git init
     docker run --name ul -it ubuntu
     apt-get update
     apt-get install -y git
     exit

     b) Save this container as an image
        docker commit ul intelliqit/ubuntu_may_25

3 Login into dockerhub
  docker login
  Enter username and password of dockerhub

4 Push the customised image
  docker push intelliqit/ubuntu_may_25

=====
Private Registry
=====
This can be created using a docker image called as "registry"
We can start this as a container and it will allow us to
push images into the registry

```

- 1 Create registry as a container
docker run --name lr -d -p 5000:5000 registry
- 2 Download an alpine image
docker pull alpine
- 3 Tag the alpine with the local registry
docker tag alpine localhost:5000/alpine
- 4 Push the image to local registry
docker push localhost:5000/alpine

```
=====
===
Day 20
=====
====
Kubernetes
=====
```

Minions: This is an individual node used in kubernetes
Combination of these minions is called as Kubernetes cluster

Master is the main machine which triggers the container orchestration
It distributes the work load to the Slaves

Slaves are the nodes that accept the work load from the master
and handle activities load balancing, autoscaling, high availability etc

```
=====
Kubernetes unmanaged setup on Centos
=====
Install, start and enable docker service

yum install -y -q yum-utils device-mapper-persistent-data lvm2 >
/dev/null 2>&1
yum-config-manager --add-repo
https://download.docker.com/linux/centos/docker-ce.repo > /dev/null 2>&1
yum install -y -q docker-ce >/dev/null 2>&1
```

```
systemctl start docker
systemctl enable docker
```

```
=====
=====
Disable SELINUX

setenforce 0
sed -i --follow-symlinks 's/^SELINUX=enforcing/SELINUX=disabled/'
/etc/sysconfig/selinux
```

```
=====
=====
Disable SWAP

sed -i '/swap/d' /etc/fstab
swapoff -a
```

```

=====
=====
Update sysctl settings for Kubernetes networking

cat >>/etc/sysctl.d/kubernetes.conf<<EOF
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
sysctl --system

=====
=====
Add Kubernetes to yum repository

cat >>/etc/yum.repos.d/kubernetes.repo<<EOF
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
        https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
EOF

=====
=====
Install Kubernetes
yum install -y kubeadm kubelet kubect1

=====
=====
Enable and start Kubernetes service

systemctl start kubelet
systemctl enable kubelet

=====
=====
Repeat the above steps on Master and slaves

=====

On Master=====
=====
Initilise the Kubernetes cluster
-----

kubeadm init --apiserver-advertise-address=ip_of_master --pod-network-
cidr=192.168.0.0/16

=====
=====

To be able to use kubect1 command to connect and interact with the
cluster,
the user needs kube config file.

mkdir /home/centos/.kube
cp /etc/kubernetes/admin.conf /home/centos/.kube/config
chown -R centos:centos /home/centos/.kube

```

```
=====
=====
```

```
Deploy calico network
kubectl create -f
https://docs.projectcalico.org/v3.9/manifests/calico.yaml
```

```
=====
=====
```

```
For slaves to join the cluster
kubeadm token create --print-join-command
```

```
=====
=====
```

```
Check the pods of kube-system are running
```

```
kubectl get pods -n kube-system
```

```
=====
===
```

```
Day 21
```

```
=====
Kubernetes
```

```
=====
```

Minions: This is an individual node used in kubernetes
Combination of these minions is called as Kubernetes cluster

Master is the main machine which triggers the container orchestration
It distributes the work load to the Slaves

Slaves are the nodes that accept the work load from the master
and handle activities load balancing, autoscaling, high availability etc

Kubernetes uses various types of Object

1 Pod: This is a layer of abstraction on top of a container. This is the smallest object that kubernetes can work on. In the Pod we have a container. The advantage of using a Pod is that kubectl commands will work on the Pod and the Pod communicates these instructions to the container. In this way we can use the same kubectl irrespective of which technology containers are in the Pod.

2 Service: This is used for port mapping and network load balancing

3 Namespace: This is used for creating partitions in the cluster. Pods running in a namespace cannot communicate with other pods running in other namespace

4 Secrets: This is used for passing encrypted data to the Pods

5 ReplicationController: This is used for managing multiple replicas of PODs and also performing scaling

6 ReplicaSet: This is similar to replicationcontroller but it is more advanced where features like selector can be implemented

7 Deployment: This used for performing all activities that a Replicaset can do it can also handle rolling update

8 Volume: Used to preserve the data even when the pods are deleted

9 Statefulsets: These are used to handle stateful application like data bases where consistency in read write operations has to be maintained.

10 Ingress: This object is used for mapping ip with domain name

Kubernetes Architecture

=====

Master Componentes

=====

Container runtime: This can be docker or anyother container technology

apiServer: Users interact with the apiServer using some cli tool like ui, command line tool like kubelet. It is the apiServer which is the gateway to the cluster It works as a gatekeeper for authentication and it validates if a specific user is having permissions to execute a specific command. Example if we want to deploy a pod or a deployment first apiServers validates if the user is authorised to perform that action and if so it passes to the next process ie the "Scheduler"

Scheduler: This process accepts the instructions from apiServer after validation and starts an application on a specific node or set of nodes. It estimates how much amount of h/w is required for an application and then checks which slave have the necessary h/w resources and instructs the kubelet to deploy the application

kubelet: This is the actual process that takes the orders from scheduler and deploy an application on a slave. This kubelet is present on both master and slave

controller manager: This check if the desired state of the cluster is always maintained. If a pod dies it recreates that pod to maintain the desired state

etcd: Here the cluster state is maintained in key value pairs. It maintains info about the slaves and the h/w resources available on the slaves and also the pods running on the slaves The scheduler and the control manager read the info from this etcd

and schedule the pods and maintain the desired state

=====

Worker components

=====

containerrun time: Docker or some other container technology

kubelet: This process interacts with container run time and the node and it start a pod with a container in it

kubeproxy: This will take the request from services to pod
It has the intellegence to forward a request to
a near by pod.Eg If an application pod wants to communicate with a db pod
then kubeproxy will take that request to the nearby pod

=====

=====

Setup of ManagedKubernetes

=====

Free

=====

1 <http://katakoda.com>

(or)

2 <http://playwithk8s.com>

Paid

=====

1 Signup for a Google cloud account

2 Click on Menu icon on top right corner--->Click on Kubernetes Engine-->Clusters

3 Click on Create cluster--->Click on Create

=====

Day 22

=====

=====

UseCase

=====

Create nginx as a pod and name it webserver

kubect1 run --image nginx webserver

To see the list of pods running

kubect1 get pods

To see more info about the pods like their ip and slave where they are running

kubect1 get pods -o wide

To delete the pod

kubect1 delete pods webserver

=====

UseCase

=====

Create mysql pod and name it mydb and go into its interactive terminal and create few tables

```
kubect1 run --image mysql:5 mydb --env MYSQL_ROOT_PASSWORD=intelliqit
```

To check the pods
kubect1 get pods

To go into the interactive terminal
kubect1 exec -it mydb -- bash

To login into the db
mysql -u root -p
Password: intellqiit

Create tables here

```
=====
Kuberentes Defintion files
=====
```

Objects in Kubernetes cluster are deployed using these defintion files
They are created using yml and they generally these 4 top level fields.

apiVersion:
kind:
metadata:
spec:

apiVersion : This specifies the code library that has to be imported to create a particualr kind of Kubernetes object

kind: Here we specify the type kubernetes object that we want to create(Pod,ReplicaSet,Deployment,Service etc)

metadata: Here we can give additional info about the Pod like the name of the Pod,some labels etc

spec: This is where exact info about the object that is created is specified like containers info port mapping,no of replicas etc

```
=====
kind                apiVersions
=====
Pod                 v1
Service             v1
Secret              v1
Namespace           v1
ReplicationController v1
Volume              v1
ReplicaSet           apps/v1
Deployment           apps/v1
StatefuleSet         apps/v1
=====
```

```
=====
Create a pod defintion file to start nginx pod with a name webserver
```

```
1 vim pod-defintion1.yml
```

```

---
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    type: proxy
    author: intelliqit
spec:
  containers:
    - name: webserver
      image: nginx

```

...

- 2 Create pod from the above file
kubect1 apply -f pod-defintion1.yml
- 3 To check the list of pods
kubect1 get pods
- 4 To delete the pods
kubect1 delete -f pod-defintion1.yml

=====

UseCase

=====

Create a postgres-pod and give the labels as author=intelliqit and type=db,also pass the necessay environment variables

```

1 vim pod-definition2.yml
apiVersion: v1
kind: Pod
metadata:
  name: postgres-pod
  labels:
    author: intelliqit
    type: db
spec:
  containers:
    - name: mydb
      image: postgres
      env:
        - name: POSTGRES_PASSWORD
          value: intelliqit
        - name: POSTGRES_USER
          value: myuser
        - name: POSTGRES_DB
          value: mydb

```

...

To create pods from the above file
kubect1 apply -f pod-defintion2.yml

=====

UseCase

=====

Create a jenkins-pod and also perfrom necessary port mapping

```

vim pod-definition2.yml
---
apiVersion: v1
kind: Pod
metadata:
  name: jenkins-pod
  labels:
    type: ci-cd
    author: intelliqit
spec:
  containers:
    - name: jenkins
      image: jenkins/jenkins
      ports:
        - containerPort: 8080
          hostPort: 8080
...

```

To create the pods from the above file
 kubectl apply -f pod-defintion3.yml

To check if the jnekins pod is running
 kubectl get pods -o wide

To accesss jenkins from browser
 kubectl get nodes -o wide
 Capture the external ip of the node where jenkins pod is running
 in browser
 externalip:8080

```

=====
===
Day 23
=====
=====
=====

```

```

ReplicationController
=====

```

This is a high level Kuberternets object that can be used for handling multiple replicas of a Pod.Here we can perfrom Load Balancing and Scalling

ReplicationController uses keys like "replicas,template" etc in the "spec" section
 In the template section we can give metadata related to the pod and also use another spec section where we can give containers information

Create a replication controller for creating 3 replicas of httpd
 vim replication-controller.yml

```

---
apiVersion: v1
kind: ReplicationController
metadata:
  name: httpd-rc
  labels:
    author: intelliqit
spec:
  replicas: 3
  template:

```

```

metadata:
  name: httpd-pod
  labels:
    author: intelliqit
spec:
  containers:
  - name: myhttpd
    image: httpd
    ports:
      - containerPort: 80
        hostPort: 8080
...

```

To create the httpd replicas from the above file
 kubectl create -f replication-controller.yml

To check if 3 pods are running and on which slaves they are running
 kubectl get pods -o wide

To delete the replicas
 kubectl delete -f replication-controller.yml

ReplicaSet

=====

This is also similar to ReplicationController but it is more advanced and it can also handle load balancing and scaling. It has an additional field in spec section called as "selector". This selector uses a child element "matchLabels" where the it will search for Pod based on a specific label name and try to add them to the cluster.

Create a replicaset file to start 4 tomcat replicas and then perform scaling

vim replica-set.yml

```

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: tomcat-rs
  labels:
    type: webserver
    author: intelliqit
spec:
  replicas: 4
  selector:
    matchLabels:
      type: webserver

```

```

template:
  metadata:
    name: tomcat-pod
    labels:
      type: webserver
  spec:
    containers:
      - name: mywebserver
        image: tomcat
        ports:

```

```
- containerPort: 8080
  hostPort: 9090
```

To create the pods from the above file
kubectl create -f replica-set.yml

Scalling can be done in 2 ways

a) Update the file and later scale it

b) Scale from the coomand prompt without updating the defintion file

a) Update the file and later scale it

Open the replicas-set.yml file and increase the replicas count from 4 to 6

```
kubectl replace -f replicas-set.yml
```

Check if 6 pods of tomcat are running

```
kubectl get pods
```

b) Scale from the coomand prompt without updating the defintion file

```
kubectl scale --replicas=2 -f replica-set.yml
```

```
=====
Deployment
=====
```

This is also a high level Kubernetes object which can be used for scalling and load balancing and it can also perfrom rolling update

Create a deployment file to run nginx:1.7.9 with 3 replicas

```
vim deployment1.yml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nginx-deployment
```

```
  labels:
```

```
    author: intelligit
```

```
    type: proxyserver
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    matchLabels:
```

```
      type: proxyserver
```

```
  template:
```

```
    metadata:
```

```
      name: nginx-pod
```

```
      labels:
```

```
        type: proxyserver
```

```
    spec:
```

```
      containers:
```

```
        - name: nginx
```

```
          image: nginx:1.7.9
```

```
          ports:
```

```
            - containerPort: 80
```

```
              hostPort: 8888
```

To create the deployment from the above file

```
kubectl create -f deployment.yml
```

To check if the deployment is running
`kubectl get deployment`

To see if all 3 pod of nginx are running
`kubectl get pod`

Check the version of nginx
`kubectl describe pods nginx-deployment | less`

```
=====
=
```

Namespace in kubernetes

```
=====
```

Namespaces are used to create partitions in the Kubernetes cluster
Pods running in different namespaces cannot communicate with
each other

To create Namespaces

```
=====
```

```
vim namespace.yml
```

```
---
```

```
apiVersion: v1
  kind: Namespace
  metadata:
    name: test-ns
...
```

```
kubectl apply -f namespace.yml
```

To see the list of namespace

```
=====
```

```
kubectl get namespace
```

Create a pod on that namespace

```
=====
```

```
vim pod-definition4.yml
```

```
---
```

```
apiVersion: v1
kind: Pod
metadata:
  name: jdk-pod
  namespace: test-ns
  labels:
    author: intelliqit
spec:
  containers:
    - name: java
      image: openjdk:12
...
```

To see list of pods in a namespace

```
=====
```

```
kubectl get pods -n test-ns
```

To delete a namespace

```
=====
```

```
kubectl delete namespace test-ns
```



```
=====
=
```

Day 24

```
=====
=
```

Kompose

```
=====
```

This is used to implement docker compose to create a multi container architecture in Kubernetes

Implementing docker compose can be done using Kompose
docker compose + docker swarm = docker stack
docker compose + Kubernetes = Kompose

Setup

```
=====
```

1 Download Kompose

```
curl -L
https://github.com/kubernetes/kompose/releases/download/v1.18.0/kompose-
linux-amd64 -o kompose
```

2 Give execute permissions

```
chmod +x kompose
```

3 Move it to PATH

```
sudo mv ./kompose /usr/local/bin/kompose
```

4 To check if the installation is successful

```
kompose version
```

Digital Ocean URL

```
=====
```

<https://www.digitalocean.com/community/tutorials/how-to-migrate-a-docker-compose-workflow-to-kubernetes>

Create a docker compose file

```
vim docker-compose.yml
```

```
---
```

```
version: '3'
```

```
services:
```

```
mydb:
```

```
image: mysql:5
```

```
environment:
```

```
MYSQL_ROOT_PASSWORD: intelliq
```

```
wordpress:
```

```
image: wordpress
```

```
ports:
```

```
- 8080:80
```

```
deploy:
```

```
replicas: 3
```

```
...
```

To setup the above architecture in Kubernetes

```
kompose up
```

To create kubernetes definition files

```
kompose convert
```

To delete the above create architecture
kompose down

Note: Practice Kompose on katokoda.com

=====

```
apiVersion: v1
kind: Pod
metadata:
  name: redis-pod
  labels:
    author: intelliqit
spec:
  containers:
    - name: redis
      image: redis
      volumeMounts:
        - name: redis-volume
          mountPath: /data/redis
  volumes:
    - name: redis-volume
      emptyDir: {}
```

Create a pod from the above file
kubectl create -f volumes.yml

To check if the volume is mounted
kubectl exec -it redis-pod -- bash

Go to the redis folder and create some files
cd redis
cat > file
Store some data in this file

To kill the redis pod install procps
apt-get update
apt-get install -y procps

Identify the process id of redis
ps aux
kill 1

Check if the redis-pod is recreated
kubectl get pods
We will see the restart count changes for this pod

If we go into this pods interactive terminal
kubectl exec -it redis-pod -- bash

We will see the data but not the s/w's (procps) we installed
cd redis
ls

ps This will not work

=====

===

Day 25

```
=====
===
=====
Service Object
=====
```

This is used for network load balancing and port mapping
It uses 3 ports
1 target port: Pod or container port
2 port: Service port
3 hostPort: Host machines port to make it accessible from external network

Service objects are classified into 3 types

- 1 clusterIP: This is the default type of service object used in Kubernetes and it is used when we want the Pods in the cluster to communicate with each other and not with external networks
- 2 nodePort: This is used if we want to access the pods from an external network and it also performs network load balancing ie even if a pod is running on a specific slave we can access it from other slave in the cluster
- 3 LoadBalancer: This is similar to Nodeport and it is used for external connectivity of a Pod and also network load balancing and it also assigns a public ip for all the slave combined together

```
=====
=====
Use Case
=====
```

Create a service definition file for port mapping an nginx pod

```
vim pod-definition1.yml
```

```
---
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    author: intellgit
    type: proxy
spec:
  containers:
    - name: appserver
      image: nginx
```

```
=====
vim service1.yml
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
```

```
    nodePort: 30008
  selector:
    author: intelliqit
    type: proxy
```

Create pods from the above pod definition file
kubectl create -f pod-definition1.yml
Create the service from the above service definition file
kubectl create -f service.yml
Now nginx can be accessed from any of the slave
kubectl get nodes -o wide
Take the external ip of any of the nodes:30008

UseCase

=====

Create a pod definition file to start a ghost pod and also create a service object of the type LoadBalancer

```
vim pod-definition7.yml
```

```
---
apiVersion: v1
kind: Pod
metadata:
  name: ghost-pod
  labels:
    author: intelliqit
    type: CMS
spec:
  containers:
  - name: ghost
    image: ghost
...
```

```
vim service2.yml
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: ghsot-service
  labels:
    author: intelliqit
spec:
  type: LoadBalancer
  ports:
  - targetPort: 2368
    port: 2368
  selector:
    type: CMS
    author: intelliqit
```

UseCase

=====

Create a pod-definition file for httpd pod and create a service object of type cluster ip for it

```
vim pod-definition8.yml
```

```
---
apiVersion: v1
```

```

kind: Service
metadata:
  name: ghsot-service
  labels:
    author: intelliqit
spec:
  type: LoadBalancer
  ports:
    - targetPort: 2368
      port: 2368
  selector:
    type: CMS
    author: intelliqit
...

```

```

vim service3.yml
---
apiVersion: v1
kind: Service
metadata:
  name: httod-service
  labels:
    author: intelliqit
spec:
  ports:
    - targetPort: 80
      port: 80
  selector:
    author: intelliqit
    type: webserver
...

```

=====

UseCase

=====

Create a deployment file of for tomcat and also create a service file of type node port

```

vim deployment3.yml
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: tomcat-deployment
  labels:
    type: appserver
spec:
  replicas: 2
  selector:
    matchLabels:
      type: appserver
  template:
    metadata:
      name: tomcat-pod
      labels:
        type: appserver
    spec:
      containers:
        - name: tomcat
          image: tomee

```

...

vim service4.yml

apiVersion: v1

kind: Service

metadata:

name: tomcat-service

labels:

author: intelliqit

spec:

type: NodePort

ports:

- targetPort: 8080

port: 8080

selector:

type: appserver

...

=====

=

=====

==

Day 26

=====

==

=====

Kubernetes Project

=====

This is a python based application which is used for accepting a vote (voting app). This application accepts the vote and passes it to a temporary db created using redis. From here the data is passed to a worker application created using .net which analyses the data and stores them permanently in a database created using postgres. From here the results can be seen on an application that is created using nodejs and this is called as resulta-app

To do this we will create 5 pod definition files and 4 service files, 2 services of type cluster ip for redis and postgres databases 2 services of type loadbalancer for python voting app and nodejs result app

Pod Definition Files

=====

vim voting-app-pod.yml

apiVersion: v1

kind: Pod

metadata:

name: voting-app-pod

labels:

name: voting-app-pod

author: intelliqit

spec:

containers:

- name: voting-app

image: dockersamples/examplevotingapp_vote

```
    ports:
      - containerPort: 80
...

```

```
vim result-app-pod.yml
---
```

```
apiVersion: v1
kind: Pod
metadata:
  name: result-app-pod
  labels:
    name: result-app-pod
    author: intelliqit
spec:
  containers:
    - name: result-app
      image: dockersamples/examplevotingapp_result
      ports:
        - containerPort: 80
...

```

```
vim worker-app-pod.yml
---
```

```
apiVersion: v1
kind: Pod
metadata:
  name: worker-app-pod
  labels:
    name: worker-app-pod
    author: intelliqit
spec:
  containers:
    - name: worker-app
      image: dockersamples/examplevotingapp_worker
...

```

```
vim redis-pod.yml
---
```

```
apiVersion: v1
kind: Pod
metadata:
  name: redis-pod
  labels:
    name: redis-pod
    author: intelliqit
spec:
  containers:
    - name: redis
      image: redis
      ports:
        - containerPort: 6379
...

```

```
vim postgres-pod.yml
---
```

```

apiVersion: v1
kind: Pod
metadata:
  name: postgres-pod
  labels:
    name: postgres-pod
    author: intelliqit
spec:
  containers:
    - name: postgres
      image: postgres
      ports:
        - containerPort: 5432
...

```

```

=====
===

```

Service Defintion file

```

=====

```

```

vim redis-service.yml

```

```

---
apiVersion: v1
kind: Service
metadata:
  name: redis-service
  labels:
    name: redis-service
    author: intelliqit
spec:
  ports:
    - port: 6379
      targetPort: 6379
  selector:
    name: redis-pod
    app: demo-voting-app
...

```

```

vim pod-service.yml

```

```

---
apiVersion: v1
kind: Service
metadata:
  name: postgres-service
  labels:
    name: postgres-service
    author: intelliqit
spec:
  ports:
    - port: 5432
      targetPort: 5432
  selector:
    name: postgres-pod
    app: demo-voting-app
...

```

Note: Since "type" is not specified in the "spec" section they will be created as clusterIP


```

vim voting-app-service.yml
---
apiVersion: v1
kind: Service
metadata:
  name: voting-app-service
  labels:
    name: voting-app-service
    author: intelliqit
spec:
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 80
  selector:
    name: voting-app-pod
    app: demo-voting-app
...

```

```

vim result-app-service.yml
---
apiVersion: v1
kind: Service
metadata:
  name: result-app-service
  labels:
    name: result-app-service
    author: intelliqit
spec:
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 80
  selector:
    name: result-app-pod
    app: demo-voting-app
...

```

The above 2 service objects are created as LoadBalancer type ie they can perform network load balancing where we can access the pod from any slave and also a single public ip will be assigned for all the slaves

=====
Day 30
=====

To deploy the above project using docker stack

```

vim voting-app.yml
---
version: '3'

services:
  voting-app:
    image: dockersamples/examplevotingapp_vote
    ports:
      - 6060:80

```

```

redis:
  image: redis
  ports:
    - 6379:6379

worker-app:
  image: dockersamples/examplevotingapp_worker

postgres:
  image: postgres
  environment:
    POSTGRES_PASSWORD: intelligit
  ports:
    - 5432:5432

result-app:
  image: dockersamples/examplevotingapp_result
  ports:
    - 7070:80

```

To deploy the above services
 docker stack deploy -c voting-app.yml my-voting-app

To see the list of nodes where the stack services are running
 docker stack ps my-voting-app

```

=====
===
To deploy the above file in kubernetes using kompose
=====
=====
Kompose
=====
This is used to implement docker compose to create a multi
container architecture in Kubernetes

```

Implementing docker compose can be done using Kompose
 docker compose + docker swarm = docker stack
 docker compose + Kubernetes = Kompose

```

Setup
=====
1 Download Kompose
  curl -L
  https://github.com/kubernetes/kompose/releases/download/v1.18.0/kompose-
  linux-amd64 -o kompose

2 Give execute permissions
  chmod +x kompose

3 Move it to PATH
  sudo mv ./kompose /usr/local/bin/kompose

4 To check if the installion is successfull
  kompose version

```

Digital Ocean URL
 =====

<https://www.digitalocean.com/community/tutorials/how-to-migrate-a-docker-compose-workflow-to-kubernetes>

```
=====
=====
Day 27
=====
=====
Stateful set
=====
---
apiVersion: v1
kind: Service
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  ports:
    - port: 80
      name: web
  clusterIP: None
  selector:
    app: nginx
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
              name: web
          volumeMounts:
            - name: www
              mountPath: /usr/share/nginx/html
  volumeClaimTemplates:
    - metadata:
        name: www
      spec:
        accessModes: [ "ReadWriteOnce" ]
        resources:
          requests:
            storage: 1Gi
```

```

=====
=====
Node affinity
=====
kubect1 get nodes --show-labels

kubect1 label nodes <your-node-name> disktype=ssd

kubect1 label nodes gke-cluster-1-default-pool-3cde7c4a-h174 disktype=ssd
=====
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: disktype
                operator: In
                values:
                  - ssd
  containers:
    - name: nginx
      image: nginx
=====

Taints and toleration
=====
Taints and Tolerations
Node affinity, is a property of Pods that attracts them to a set of nodes
(either as a preference or a hard requirement). Taints are the opposite -
- they allow a node to repel a set of pods.

Tolerations are applied to pods, and allow (but do not require) the pods
to schedule onto nodes with matching taints.

Taints and tolerations work together to ensure that pods are not
scheduled onto inappropriate nodes. One or more taints are applied to a
node; this marks that the node should not accept any pods that do not
tolerate the taints.

To create a taint for a node
kubect1 taint nodes node1 node=intelliqit:NoSchedule

To delete the taint
kubect1 taint nodes node1 node=intelliqit:NoSchedule-

Pod definition file to use the above taint
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    author: intelliqit
spec:
  containers:
    - name: mynginx

```

```
    image: nginx
tolerations:
- key: "node"
  operator: "Equal"
  value: "intelliqit"
  effect: "NoSchedule"
```

=====

Helm and Kubernetes

Helm is a package manager. Package managers automate the process of installing, configuring, upgrading, and removing computer programs

Helm has two elements, a client (Helm) and a server (Tiller). The server element runs inside a Kubernetes cluster and manages the installation of charts.